

Unsafe Rust

CIS 198 Lecture 12

Unsafe Rust

- Rust's safety checks are convenient, but highly conservative.
- Some Rust programs are completely safe, but the compiler won't accept them.
- Rust has "unsafe" parts to allow you to escape some of the compiler's restrictions.
 - Sometimes you need to "just [frob](#) some dang bits." -[Gankro](#)

Unsafe Rust

- Unsafe Rust is useful for:
 - Talking to other languages (e.g. C)
 - Writing some low-level abstraction that `std` doesn't provide
 - Using particular types that convey no ownership semantics
 - Writing `std` itself!

Unsafe Rust

- All unsafe code must be delimited with the **unsafe** keyword.
- Unsafe code can be wrapped in an **unsafe** block.
- Entire functions can be marked as unsafe to call.
 - Code which calls unsafe functions must be in **unsafe** blocks.

```
unsafe fn foo() { }

fn main() {
    unsafe {
        foo();
    }
}
```

- Traits can be designated unsafe.
 - Thus, their **impls** must be also marked as unsafe.

```
unsafe trait Sync { }
unsafe impl Sync for Foo { }
```

Unsafe vs. Safe

- "Safe" operations are ones that don't break any of Rust's ownership, type-checking, or memory safety rules.
- These are dangerous, but are still considered "safe" in Rust:
 - Deadlocks
 - Memory leaks
 - Integer overflow
 - Exiting without running destructors

Unsafe vs. Safe

- "Unsafe" operations are ones that *might* have dangerous effects that the Rust compiler cannot guard against by analysis.
- These are bad, and are only possible in **unsafe** Rust:
 - Data races
 - Dereferencing invalid pointers
 - Reading uninitialized memory
 - Creating invalid primitives
 - etc.

Unsafe Rust

- What does **unsafe** let you do?
 - Dereference raw pointers
 - Mutate **static** variables
 - Implement **unsafe** traits
 - Call **unsafe** functions
- That's it!
- **unsafe** does not disable the borrow checker or otherwise let you do anything that would break regular Rust semantics.
- Everything that's normally safe is still safe in an **unsafe** context.
- Everything else that the compiler would reject is still disallowed.

Where Is **unsafe** Used?

- All over the place!
- Many, many types in **std** are built on **unsafe** in some way.
 - **Vec**
 - **Cell/RefCell**
 - **Box**
 - **Rc/Arc**
 - **Send/Sync**
 - **HashMap**
 - Lots and lots of other types

Correctness of Unsafe Code

- If `std` is built on `unsafe`, how can it possibly be correct?
- `unsafe` is used in very careful ways within `std`.
- There's no formalization of `std`'s safety, so you pretty much need to just trust that `std` is implemented correctly.
 - `Rustbelt` is an ongoing project to formally verify the encapsulation of unsafe code. Seeking PhD students!

Correctness of Unsafe Code

- As it turns out, using **unsafe** pollutes the safety of everything that can directly modify data that the **unsafe** code depends on.
 - And if not written correctly, safe functions which use **unsafe** will pollute the safety of any code which uses it.
- For example, take this (simplified) definition of **Vec**.
 - **ptr** is a pointer to the data the **Vec** owns.
 - **cap** is the **Vec**'s capacity, **len** is its length.
 - **len** should never be greater than **cap**, but this is not enforced.

```
pub struct Vec<T> {  
    ptr: *mut T,  
    cap: usize,  
    len: usize,  
}
```

Correctness of Unsafe Code

- Now consider this new method on **Vec**, written in totally safe code.
- Despite this code's total safety, it does a dangerous thing.
 - The length of the **Vec** is now potentially longer than its capacity.
- This method allows clients of **Vec** to access uninitialized memory!

```
impl<T> Vec<T> {  
    // ...  
    pub fn evil(&mut self) {  
        self.len += 2;  
    }  
}
```

¹ Example taken from [The Scope of Unsafe](#)

Borrow Splitting

- Mutual exclusivity of mutable references is safe, but can be very limiting.
- Struct borrows can be split to allow multiple mutable borrows, as the borrow checker understands structs well enough:

```
struct Point {  
    x: i32,  
    y: i32,  
}  
  
let mut p = Point { x: 0, y: 1 };  
let x1 = &mut p.x;  
let y1 = &mut p.y;  
*x1 += 1;  
*y1 += 1; // OK
```

Borrow Splitting

- However, arrays and slices cause problems.
- While the borrow checker *could* understand this, more complex examples with trees, etc., are unfeasible.

```
let mut arr = [1, 2, 3];  
let x = &mut arr[0];  
let y = &mut arr[1];  
// ^ cannot borrow `arr[..]` as mutable more than once at a time  
let z = &arr[2];
```

Borrow Splitting

- Unsafety can allow us to carefully overcome this limitation.
- Slices can be split into parts and subsequently lent out.
- `std` provides a function `split_at_mut`, that looks like this:

```
fn split_at_mut(&mut self, mid: usize) -> (&mut [T], &mut [T]) {  
    let len = self.len();  
    let ptr = self.as_mut_ptr();  
    assert!(mid <= len);  
    unsafe {  
        (from_raw_parts_mut(ptr, mid),  
         from_raw_parts_mut(ptr.offset(mid as isize), len - mid))  
    }  
}
```

- This is a classic example of how to use unsafe.
 - We know this function works safely, but the borrow checker can't validate this property.
 - Unsafe code can be used as a careful escape hatch.

Raw Pointers

- Rust defines two types of C-like "raw pointers":
 - `*const T`, an immutable pointer to a `T`
 - `*mut T`, a mutable pointer to a `T`
- Raw pointers are unsafe to dereference.
- Raw pointers confer no ownership semantics, just like C pointers.
- Raw pointers may be initialized to null with `null()` and `null_mut()`, respectively.
 - These are actually safe - only dereferencing is unsafe.

Raw Pointers

- Raw pointers may be (safely) created in several ways:

```
// Reference coercion
let mut x = 0i32;
let const_ptr = &x as *const i32;
let mut_ptr = &mut x as *mut i32; // Allowed to alias pointers

let box_y = Box::new(1);
let raw_y = &*box_y as *const i32; // Does not consume `box_y`
```


Raw Pointers

- Raw pointers may be (safely) created in several ways:

```
// Box consumption
let raw_y_2 = Box::into_raw(box_y);

// ... to properly clean up, later ...
unsafe { drop(Box::from_raw(raw_y_2)); }
```

- Raw pointers may need to be deallocated manually.
 - If you've taken ownership of the pointer, not if you've aliased it.
- Raw pointers *can*, but should not, be used after their lifetime expires.
 - Doing so will cause a segmentation fault.

Unique

- A wrapper around `*mut T` to indicate that the `Unique` struct *owns* the pointer.
- Confers regular Rust ownership semantics, unlike `*mut T`.
- Implies `T` should not be modified without a unique path to `T`.
- Unsafe to create, dereference, etc.
- Useful for building abstractions such as `Box`.
- Currently unstable :(
 - Available on nightly behind a feature gate:

```
#![feature(unique)]
use std::ptr::Unique;

let mut x = 0i32;
let mut_ptr = &mut x as *mut i32;
unsafe {
    let unique_ptr = Unique::new(mut_ptr);
}
```

Shared

- A wrapper around `*mut T` to indicate that the `Shared` struct has *shared ownership* of the pointer.
- Confers ownership semantics, unlike `*mut T`.
- Unsafe to create, dereference, etc.
- Useful for building abstractions such as `Rc`, `Arc`.
- Also currently unstable :(

UnsafeCell

- A wrapper around a **T** that provides interior mutability.
- Implies that **T** may be modified using **unsafe** operations.
- A primitive cell type that is the basis for **Cell/RefCell**.

Uninitialized Memory

- If you really, *really* want some uninitialized memory, you can make some with `std::mem::uninitialized<T>()`.
 - This function just pretends to make a value of type `T`.
 - This function is unsafe to invoke.
- **Reading uninitialized memory always has undefined behaviour!**
- Writing to uninitialized memory may also be undefined.
 - The compiler believes the value is initialized, so it may try to drop a value that isn't there, causing a panic.
- The only safe way to initialize uninitialized memory is by using one of the functions from `std::ptr`.
 - `write`, `copy`, or `copy_nonoverlapping`

Leaking Memory

- Leaking memory is a safe operation, and can be done with `std::mem::forget()`.
- You rarely ever want to do this, but you might want to if:
 - you have an uninitialized value that you don't want to get dropped
 - you have two copies of a value, but only want to drop one to avoid a double free
 - you want to transfer a resource across the FFI boundary and into another language
- This function is safe because Rust does not guarantee that destructors are always run.

Transmutation

- Is the type system just too... pesky for you?
- Want to turn any type into any other type as much as you please?
- `mem::transmute` may be right for you!¹

¹Ask your compiler before using `mem::transmute`. Side effects may include your laundry being eaten by nasal goblins.

Transmutation

- `mem::transmute<T, U>` takes a value of type `T` and reinterprets it to be of type `U`.
- The only restriction on this function is that `T` and `U` must be the same size.
- Arbitrary type transmutation can have wildly undefined behaviour!

Transmutation

- Creating an instance of any type with invalid state causes arbitrary behaviour.
 - e.g. a **Vec** with a capacity smaller than its length.
 - The above may or may not be possible with an actual **Vec**, it's just an example.
- **transmute** can create invalid primitives, e.g. a **bool** that is neither **0** nor **1** underneath.
- **transmute** can turn an **&T** into an **&mut T**, which is *wildly, always* undefined
- **transmute** on a reference without an explicit lifetime produces an *unbounded* lifetime.

Transmutation

- There's also a variant named `transmute_copy` which is *even more unsafe!*
- `transmute_copy<T, U>` copies `sizeof<U>` bytes from an `&T` and interprets them as a `U`.
- No more pesky size check like in regular `transmute`.
- `U` may be larger than `T`, which causes undefined behaviour.

Unbounded Lifetimes

- Unsafe code often produces a reference or lifetime out of nowhere.
 - These lifetimes are unbounded.
- Commonly, these are produced from dereferencing a raw pointer.
- Unbounded lifetimes expand to become as large as they need to be.
- This may produce lifetimes more powerful than 'static.
- Generally, you can think of these as 'static.
 - Almost no references are actually 'static, though.

Implementing **Vec**

- Let's take a deep dive into how `std::vec::Vec` is actually implemented.
- This code will only compile on nightly, since it uses unstable features.
- Warning: this is going to get pretty advanced.

¹All content taken from [The Rustonomicon](#)

Vec Layout

- A **Vec** has three parts:
 - a pointer to the allocated data
 - the size of the allocation
 - the number of elements in the vector
- Naively, this translates into this struct:

```
pub struct Vec<T> {  
    ptr: *mut T,  
    cap: usize,  
    len: usize,  
}
```

- Simple, right?

Vec Layout


- Sadly, this won't quite work.
- Some lifetime variance problems:
 - An `&Vec<&'static str>` can't be used in place of an `&Vec<&'a str>`.
- Some drop checker problems:
 - `*mut T` conveys no ownership.
 - The drop checker assumes we own no values of type `T`.

Vec Layout

- Using a **Unique** in place of an ***mut T** solves these problems!
 - **Unique** conveys ownership.
 - It also lets us be **Send** & **Sync** if **T** is, and auto-derefs to ***mut T**.
 - It also is guaranteed to never be null, allowing null pointer optimizations.

```
#![feature(unique)]  
use std::ptr::{Unique, self};  
  
pub struct Vec<T> {  
    ptr: Unique<T>,  
    cap: usize,  
    len: usize,  
}
```

Vec Allocation

- Now that we use a **Unique** to store our data, what does an empty **Vec** look like?
 - It can't have a null pointer...
 - ...but we don't want to allocate any data!
- Turns out, we can just fill the **Vec** with garbage! 
- **cap** is 0, so we don't need to worry about accidentally accessing garbage.
- **std** exposes a value as **alloc::heap::EMPTY** to represent this value.

Vec Allocation

```
#![feature(alloc, heap_api)]
use alloc::heap::EMPTY;

impl<T> Vec<T> {
    fn new() -> Self {
        unsafe {
            assert!(mem::size_of::<T>() != 0);
            Vec { ptr: Unique::new(heap::EMPTY as *mut _),
                  len: 0,
                  cap: 0 }
        }
    }
}
```

Vec Allocation

- Now that we can allocate no space, we need to figure out how to allocate space.
- The **heap** module is our friend here.
 - This lets us talk to Rust's allocator (jemalloc by default).
- We also need to figure out how to handle an out-of-memory condition.
 - **panic!** is no good, since it can cause allocations.
 - **std** usually executes an illegal instruction to crash the program.
- We'll define an out-of-memory error like so:

```
fn oom() { ::std::process::exit(-9999); }
```

Vec Allocation

- Our logic for growing the **Vec** is roughly this:

```
if cap == 0 {  
    allocate()  
    cap = 1  
} else {  
    reallocate()  
    cap *= 2  
}
```

- Unfortunately, this is not so easy in practice...

Vec Allocation

- Because of some serious LLVM shenanigans, a `Vec` can only contain `isize::MAX` elements.
 - This *also* means we only care about byte-sized allocations, since e.g. `isize::MAX u16s` will only just fit in memory.
- Zero-size types are also tricky to allocate due to LLVM.
- tl;dr Read the ch. 10 in the Rustonomicon, this info is totally out of scope.

Vec Allocation

- Time to actually allocate data.
- The **heap** module exposes an **allocate** function that puts data into the heap.
- It also exposes a **reallocate** function that takes an existing allocation and resizes it.
- Our pseudocode from a few slides ago can now be implemented, with a few extra parts.

Vec Allocation

```
fn grow(&mut self) {
    unsafe {
        let align = mem::align_of::<T>();
        let elem_size = mem::size_of::<T>();
        let (new_cap, ptr) = if self.cap == 0 {
            let ptr = heap::allocate(elem_size, align);
            (1, ptr)
        } else {
            let new_cap = self.cap * 2;
            let old_num_bytes = self.cap * elem_size;
            assert!(old_num_bytes <=
                (::std::isize::MAX as usize) / 2);
            let new_num_bytes = old_num_bytes * 2;
            let ptr = heap::reallocate(*self.ptr as *mut _,
                old_num_bytes, new_num_bytes, align);
            (new_cap, ptr)
        }
    };

    if ptr.is_null() { oom(); }
    self.ptr = Unique::new(ptr as *mut _);
    self.cap = new_cap;
}
```

Vec Push & Pop

- Actual functionality! Woo!
- **push** only needs to check if the **Vec** is full to grow, write to the next available index, and increment the length.
 - Be careful! This method shouldn't read from the memory it writes to, since this is uninitialized.
 - Even calling **v[idx] = x** will try to drop the old value of **v[idx]**.
- To avoid reading uninitialized memory, use **ptr::write**.

```
pub fn push(&mut self, elem: T) {  
    if self.len == self.cap { self.grow(); }  
    unsafe {  
        ptr::write(self.ptr.offset(self.len as isize), elem);  
    }  
  
    self.len += 1;  
}
```

Vec Push & Pop

- With `pop`, the data being removed is initialized, so we can read it.
- Unfortunately, just moving the value out doesn't work.
 - This would leave memory uninitialized.
- `ptr::read` does the trick here, copying the data out.

```
fn pop(&mut self) -> Option<T> {  
    if self.len == 0 {  
        None  
    } else {  
        self.len -= 1;  
        unsafe {  
            Some(ptr::read(self.ptr.offset(self.len as isize)))  
        }  
    }  
}
```


Vec Deallocation

- Now we want to be able to deallocate a **Vec**.
- This actually requires a manual implementation of **Drop**.
- The whole **Vec** can be deallocated by popping it down to zero, then calling **heap::deallocate**.
- Obviously, we shouldn't try to deallocate anything if the **Vec** has no memory allocated.

Vec Deallocation

```
impl<T> Drop for Vec<T> {  
    fn drop(&mut self) {  
        if self.cap != 0 {  
            while let Some(_) = self.pop {}  
            let align = mem::align_of::<T>();  
            let elem_size = mem::size_of::<T>();  
            let num_bytes = elem_size * self.cap;  
            unsafe {  
                heap::deallocate(*self.ptr as *mut _, num_bytes,  
                                align);  
            }  
        }  
    }  
}
```

Vec Deref

- How do you implement `Deref` & `DerefMut` for `Vec`?
- For `Vec<T>`, these methods return an `&[T]`/`&mut [T]`, respectively.
- Basically, this boils down to calling `::std::slice::from_raw_parts(*self.ptr, self.len)`.

Vec Insert & Remove

- Inserting into a **Vec** needs to shift all elements over to the right from the given index.
- Fortunately, we have an easy way to do this using **ptr::copy**, which is like **memmove** from C.
- **ptr::copy** copies some chunk of memory from here to there, and the source and destination may overlap.
- Inserting at index **i** shifts **[i .. len]** to **[i+1 .. len+1]**, using the old **len**.

Vec Insert & Remove

```
pub fn insert(&mut self, index: usize, elem: T) {  
    assert!(index <= self.len);  
    if self.cap == self.len { self.grow(); }  
  
    unsafe {  
        if index < self.len {  
            ptr::copy(self.ptr.offset(index as isize),  
                      self.ptr.offset(index as isize + 1),  
                      self.len - index);  
        }  
        ptr::write(self.ptr.offset(index as isize), elem);  
        self.len += 1;  
    }  
}
```

Vec Insert & Remove

- Removing elements by index just behaves in the opposite direction.
- Shift all elements from `[i+1 .. len+1]` to `[i .. len]` using the *new* `len`.

Vec Insert & Remove

```
pub fn remove(&mut self, index: usize) -> T {  
    assert!(index < self.len);  
    unsafe {  
        self.len -= 1;  
        let t = ptr::read(self.ptr.offset(index as isize));  
        ptr::copy(self.ptr.offset(index as isize + 1),  
                  self.ptr.offset(index as isize),  
                  self.len - index);  
        t  
    }  
}
```

Rust FFI

- FFI: Foreign Function Interface.
- FFI means calling one language from another.
- For us, this means:
 - Rust calling C
 - C calling Rust
 - Other languages (e.g. Python/Ruby) calling Rust
- Why C? C is the de facto language used for FFI.
 - Often referred to as the programming *lingua franca*.

Calling C from Rust

- Why?
 - Sometimes you need to call another library - e.g. OpenSSL - that would be too costly to reimplement in Rust.
 - Sometimes you need to interface with a particular language; many native programming languages can make C bindings.
- Calling foreign functions from Rust is **unsafe**!
 - Because, of course, C is unsafe.

Calling C from Rust

- Compile C to static libraries (.a/.lib).
 - `cc -c -o foo.o foo.c`
 - `ar rcs libfoo.a foo.o`
- Or to dynamic libraries (.so/.dylib/.dll).
 - `cc -c -fPIC -o foo.o foo.c`
 - `cc -shared -fPIC -o libfoo.so foo.o`

Calling C from Rust

- In C:

```
int32_t foo() { return 10; }
```

- In Rust:

```
#[link(name = "foo", kind = "static")] // links libfoo.a.  
extern {  
    fn foo() -> i32;  
}  
  
#[link(name = "foo")] // links libfoo.so.  
extern { // By default, this is also  
    fn foo() -> i32; // statically linked.  
}
```

- Calling foreign functions is unsafe:

```
fn main() {  
    println!("foo: {}", unsafe { foo() });  
}
```

Calling C from Rust

- In C APIs, it's common to use incomplete struct definitions to define types whose implementation shouldn't be exposed to users.
 - These can't be instantiated by the user.

```
struct OpaqueThing;
```

- To represent a type like this in Rust (for FFI), we want to make a type that we can't instantiate! So we do this:

```
enum OpaqueThing { }
```

- Now, we can't have one of these, but we *can* have a pointer to one: `*mut OpaqueThing`.
 - Pointers to opaque types are very common in C interfaces.

Calling Rust from C

- Why?
 - Writing particularly dubiously-safe code in Rust.
 - Writing any part of a C program in a nicer language.

Rust from C: `#[repr(C)]`, `extern "C"`

- Rust has its own rules about memory layout and calling convention, which are different from C's.
- In order to call from C (or any other language!), we have to use C rules. (C doesn't have generics, enums with fields, etc.)

```
#[repr(C)]
pub enum Color { Red = 1, Blue, Green = 5, Yellow }

#[repr(C)]
pub struct Bikeshed { height: f64, area: f64 }

extern "C" pub fn paint(bs: Bikeshed, c: Color) { /* ... */ }
```

- Use opaque structs to hide stuff that C can't use:

```
struct X<T>(T); // C doesn't have generics.

#[repr(C)]
pub struct Xi32(X<i32>); // This struct hides the type parameter.
```

Rust from C: Static Linking

- Compile Rust to static libraries (.a) and link at build time.
 - Compile to `target/release/libfoo.a`.
- `Cargo.toml`:

```
[lib]
crate-type = ["staticlib"]
```

- `$ cc -Ltarget/release -lfoo -o main main.c`

```
#include <stdint.h>

int32_t foo(); // Function prototype for Rust function

int main() {
    printf("foo() -> %d\n", foo());
}
```

Rust from C: Dynamic Linking

- Compile Rust to dynamic libraries (**.so**) and load at runtime.
 - Compile to **target/release/libfoo.so**.
- **Cargo.toml**:

```
[lib]
crate-type = ["dylib"]
```

- In C:

```
#include <dlfcn.h>

int main() {
    void *handle = dlopen("target/release/libfoo.so", RTLD_LAZY);
    int32_t (*foo)() = dlsym(handle, "foo"); // get function ptr
    // plus error checking

    printf("foo() -> %d\n", foo());

    dlclose(handle);
}
```


Calling Rust from (e.g.) Python

- Why?
 - Many languages (Python, Ruby, etc.) are extremely slow compared with C and Rust.
 - Traditionally, high-performance functions are written in C or C++, and called from scripting languages.
 - Rust is an ideal alternative high-performance language.
 - Great for data processing, scientific computing, multithreaded code, etc.

Rust from Python

- Compile Rust to a dynamic library (.so/.dylib/.dll).
- In Python:

```
import ctypes
libfoo = ctypes.CDLL("target/release/libfoo.so")
print("foo() -> {}".format(libfoo.foo()))
```

Rust from Ruby

- Compile Rust to a dynamic library (.so/.dylib/.dll).
- In Ruby:

```
require 'fiddle'
require 'fiddle/import'

module Foo
  extend Fiddle::Importer
  dlload 'target/release/libfoo.so'
  extern 'int foo()'
end

puts Foo.foo()
```

rust-bindgen

- Generates Rust bindings for a C header.
 - Install via `cargo install bindgen`

```
typedef char my_bool;  
typedef struct st_mysql_field { char *name; /* */ } MYSQL_FIELD;  
my_bool STDCALL mysql_thread_init(void);
```

- `$ bindgen -l mysql -match mysql.h -o mysql.rs /usr/include/mysql/mysql.h`

```
pub type my_bool = ::std::os::raw::c_char;  
  
#[repr(C)] #[derive(Copy)] pub struct Struct_st_mysql_field {  
    pub name: *mut ::std::os::raw::c_char, /* */ }  
impl ::std::clone::Clone for Struct_st_mysql_field { /* */ }  
impl ::std::default::Default for Struct_st_mysql_field { /* */ }  
pub type MYSQL_FIELD = Struct_st_mysql_field;  
  
#[link(name = "mysql")]  
extern "C" { pub fn mysql_thread_init() -> my_bool; }
```

rusty-cheddar

- Generates C bindings (header) for Rust code.
 - Generation is done in the `build.rs` script.
- Only C-like enums (no fields).
- No diverging (`-> !`) functions.

¹ disclaimer: we haven't tried this.